



COMP 412  
FALL 2010

# *Code Optimization, Part II*

## *Regional Techniques*

### *Comp 412*

Copyright 2010, Keith D. Cooper & Linda Torczon, all rights reserved.

Students enrolled in Comp 412 at Rice University have explicit permission to make copies of these materials for their personal use.

Faculty from other educational institutions may use these materials for nonprofit educational purposes, provided this copyright notice is preserved.



# Last Lecture

---

## Introduced concept of a redundant expression

*An expression,  $x+y$ , is redundant at point  $p$  if, along each path from the procedure's entry point to  $p$ ,  $x+y$  has already been evaluated and neither  $x$  nor  $y$  has been redefined.*

- If  $x+y$  is redundant at  $p$ , we can save the results of those earlier evaluations and reuse them at  $p$ , avoiding evaluation
- In a single block, we need only consider one such path
- We developed an algorithm for redundancy elimination in a single basic block
- Two pieces to the problem
  - Proving that  $x+y$  is redundant
  - Rewriting the code to eliminate the redundant evaluation
- Value numbering does both for straightline code

# Local Value Numbering

(Recap)



## The LVN Algorithm, with bells & whistles

for  $i \leftarrow 0$  to  $n-1$

1. get the value numbers  $V_1$  and  $V_2$  for  $L_i$  and  $R_i$

2. if  $L_i$  and  $R_i$  are both constant then

    evaluate  $L_i \text{ Op}_i R_i$ , assign it to  $T_i$ , and mark  $T_i$  as a constant

3. if  $L_i \text{ Op}_i R_i$  matches an identity then

    replace it with a copy operation or an assignment

4. if  $\text{Op}_i$  commutes and  $V_1 > V_2$  then

    swap  $V_1$  and  $V_2$

5. construct a hash key  $\langle V_1, \text{Op}_i, V_2 \rangle$

6. if the hash key is already present in the table then

    replace operation  $I$  with a copy into  $T_i$  and mark  $T_i$  with the VN

    else

        insert a new VN into table for hash key & mark  $T_i$  with the VN

Block is a sequence of  $n$  operations of the form

$$T_i \leftarrow L_i \text{ Op}_i R_i$$

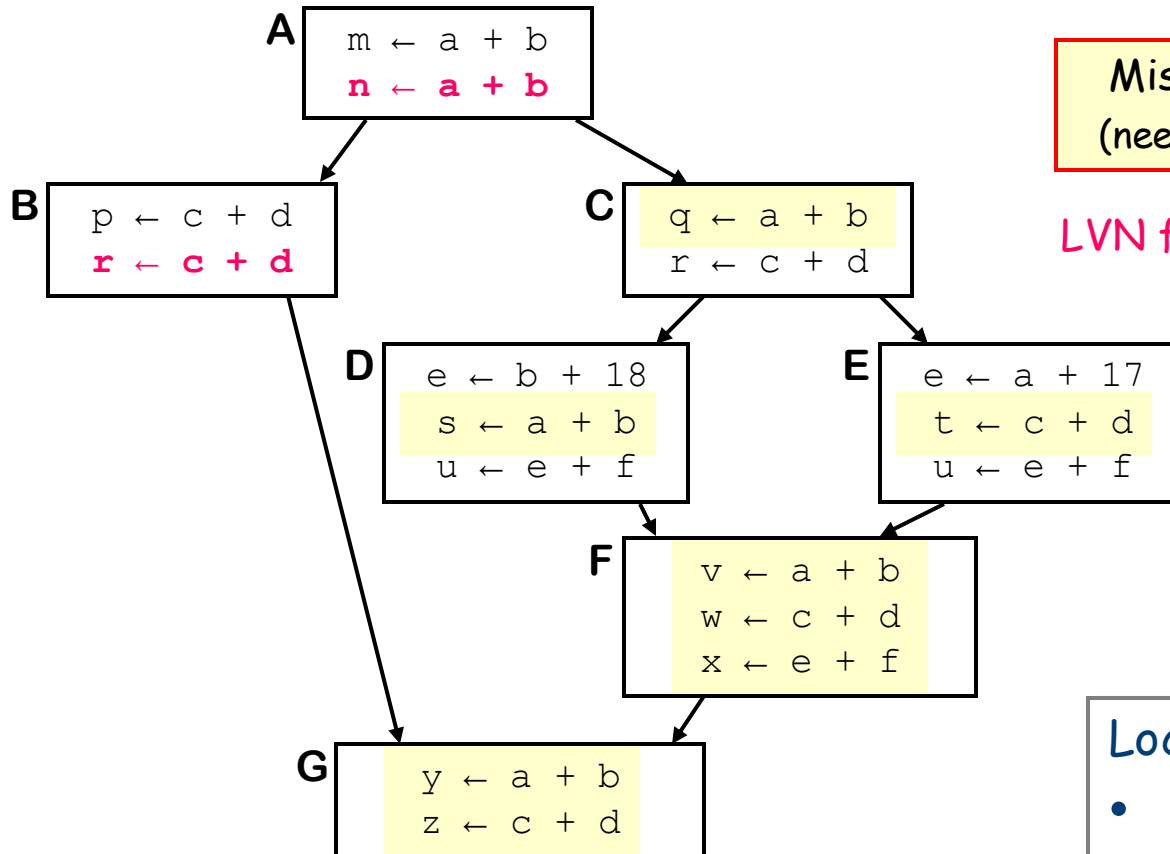
Constant folding

Algebraic identities

Commutativity



# Local Value Numbering



Missed opportunities  
(need stronger methods)

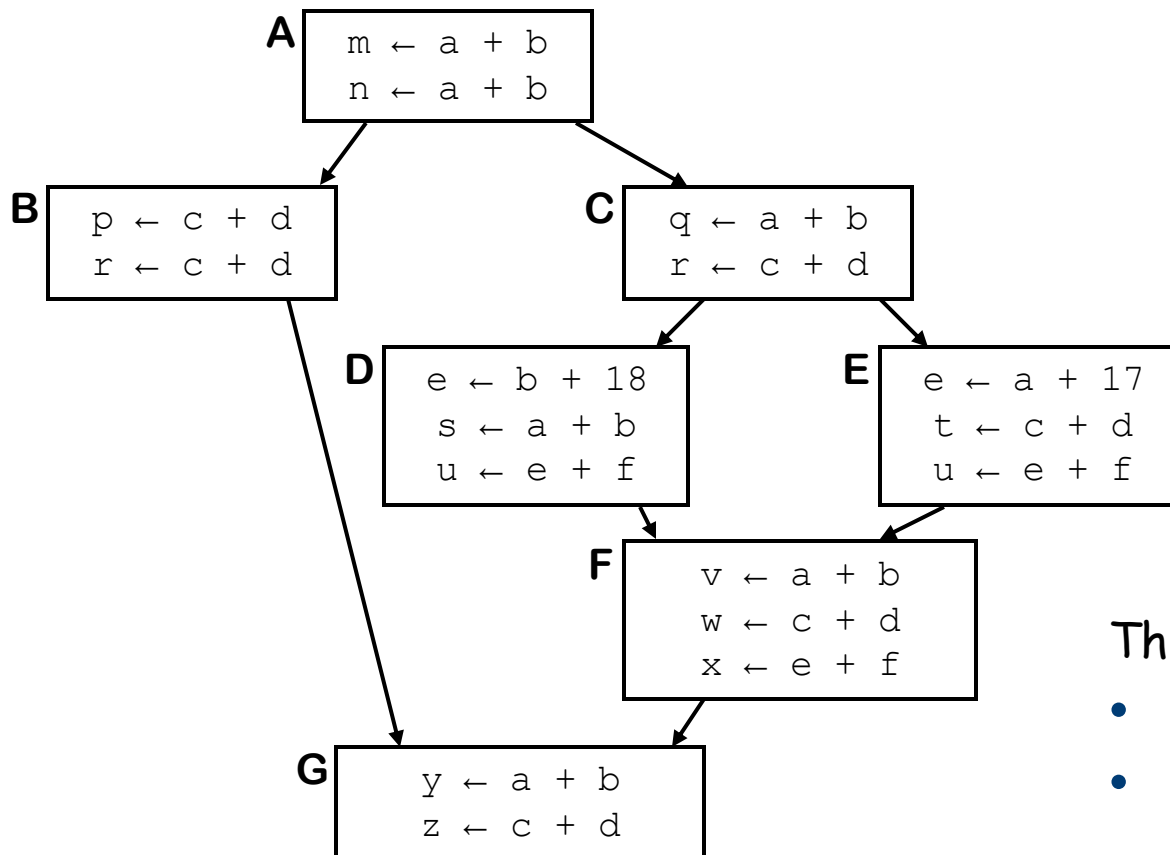
LVN finds these redundant ops

## Local Value Numbering

- 1 block at a time
- Strong local results
- No cross-block effects



# Terminology



## Control-flow graph (CFG)

- Nodes for basic blocks
- Edges for branches
- Basis for much of program analysis & transformation

This CFG,  $G = (N, E)$

- $N = \{A, B, C, D, E, F, G\}$
- $E = \{(A, B), (A, C), (B, G), (C, D), (C, E), (D, F), (E, F), (F, G)\}$
- $|N| = 7, |E| = 8$



# Scope of Optimization

---

In scanning and parsing, "scope" refers to a region of the code that corresponds to a distinct name space.

In optimization "scope" refers to a region of the code that is subject to analysis and transformation.

- Notions are somewhat related
- Connection is not necessarily intuitive

Different scopes introduces different challenges & different opportunities

Historically, optimization has been performed at several distinct scopes.



# Scope of Optimization

---

A basic block is a maximal length sequence of straightline code.

## Local optimization

- Operates entirely within a single basic block
- Properties of block lead to strong optimizations

## Regional optimization

- Operate on a region in the CFG that contains multiple blocks
- Loops, trees, paths, extended basic blocks

## Whole procedure optimization (*intraprocedural*)

- Operate on entire CFG for a procedure
- Presence of cyclic paths forces analysis then transformation

## Whole program optimization (*interprocedural*)

- Operate on some or all of the call graph (*multiple procedures*)
- Must contend with call/return & parameter binding



# A Comp 412 Fairy Tale

---

We would like to believe optimization developed in an orderly fashion

- Local methods led to regional methods
- Regional methods led to global methods
- Global methods led to interprocedural methods

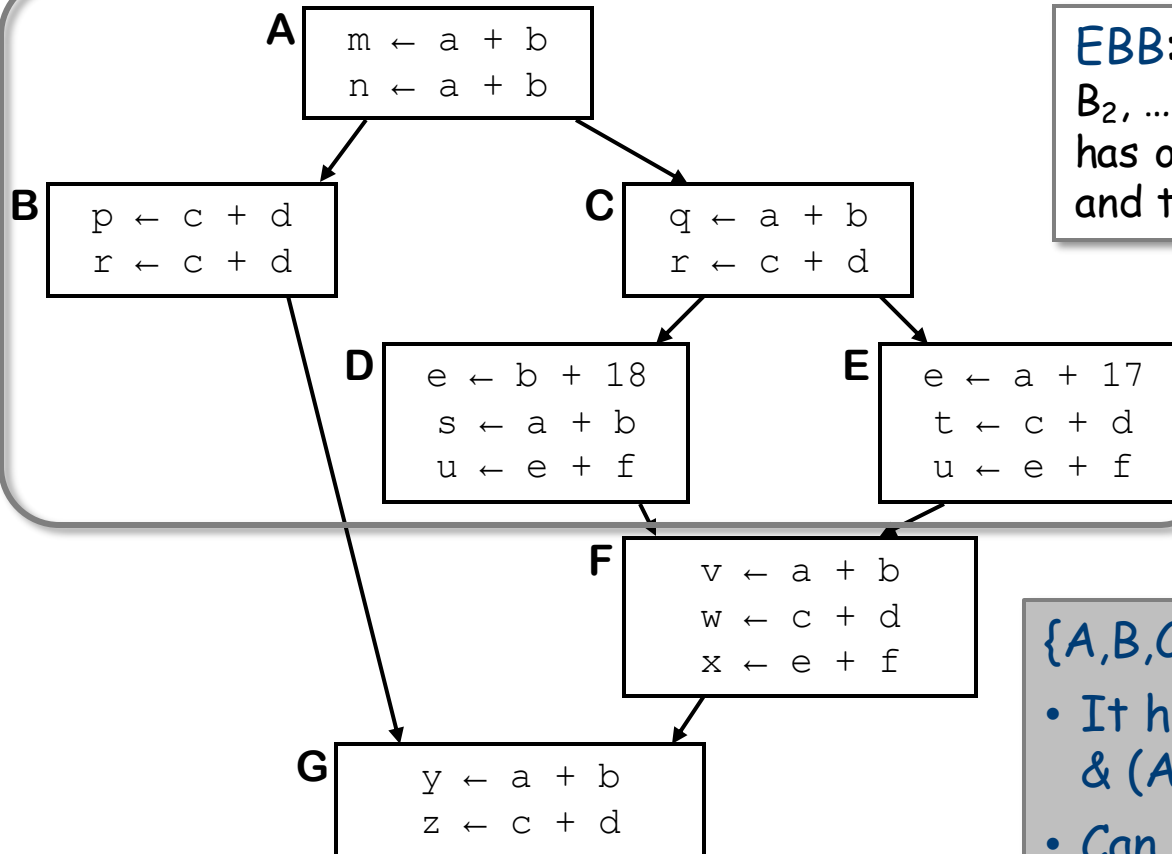
It did not happen that way

- First compiler, FORTRAN, used both local & global methods
- Development has been scattershot & concurrent
- Scope appears to relate to the inefficiency being attacked, rather than the refinement of the inventor.





# Superlocal Value Numbering



**EBB:** A maximal set of blocks  $B_1, B_2, \dots, B_n$  where each  $B_i$ , except  $B_1$ , has only exactly one predecessor and that block is in the EBB.

$\{A, B, C, D, E\}$  is an EBB

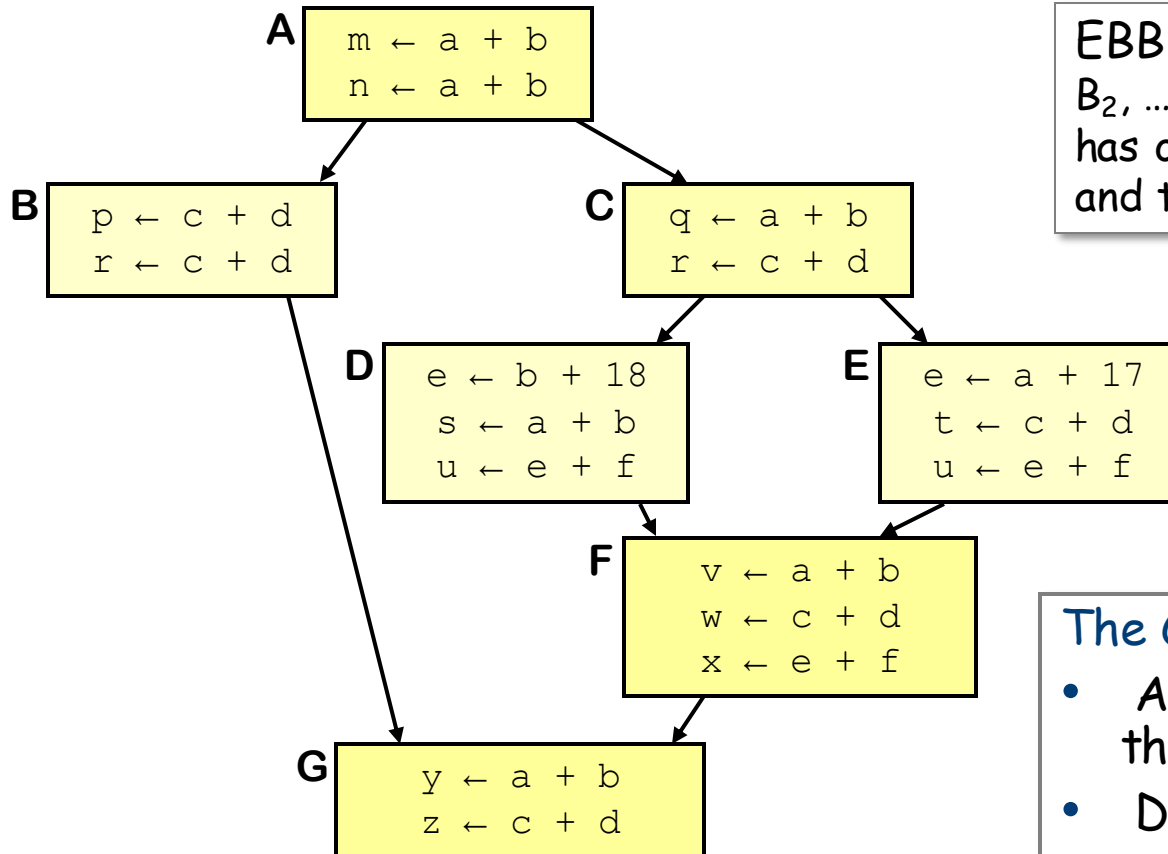
- It has 3 paths:  $(A, B)$ ,  $(A, C, D)$ , &  $(A, C, E)$
- Can sometimes treat each path as if it were a block

$\{F\}$  &  $\{G\}$  are degenerate EBBs

Superlocal: "applied to an EBB"



# Superlocal Value Numbering



EBB: A maximal set of blocks  $B_1, B_2, \dots, B_n$  where each  $B_i$ , except  $B_1$ , has only exactly one predecessor and that block is in the EBB.

## The Concept

- Apply local method to paths through the EBBs
- Do  $\{A, B\}$ ,  $\{A, C, D\}$ , &  $\{A, C, E\}$
- Obtain reuse from ancestors
- Avoid re-analyzing A & C
- Does not help with F or G

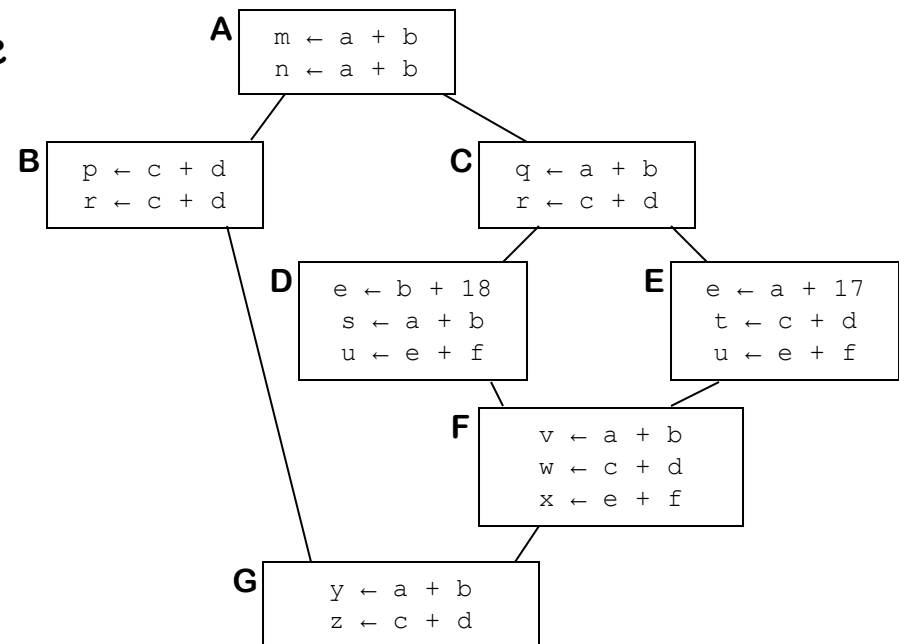


# Superlocal Value Numbering

## Efficiency

- Use A's table to initialize tables for B & C
- To avoid duplication, use a scoped hash table
  - A, AB, A, AC, ACD, AC, ACE, F, G
- Need a VN  $\rightarrow$  name mapping to handle kills
  - Must restore map with scope
  - Adds complication, not cost

"kill" is a re-definition of some name





# Superlocal Value Numbering

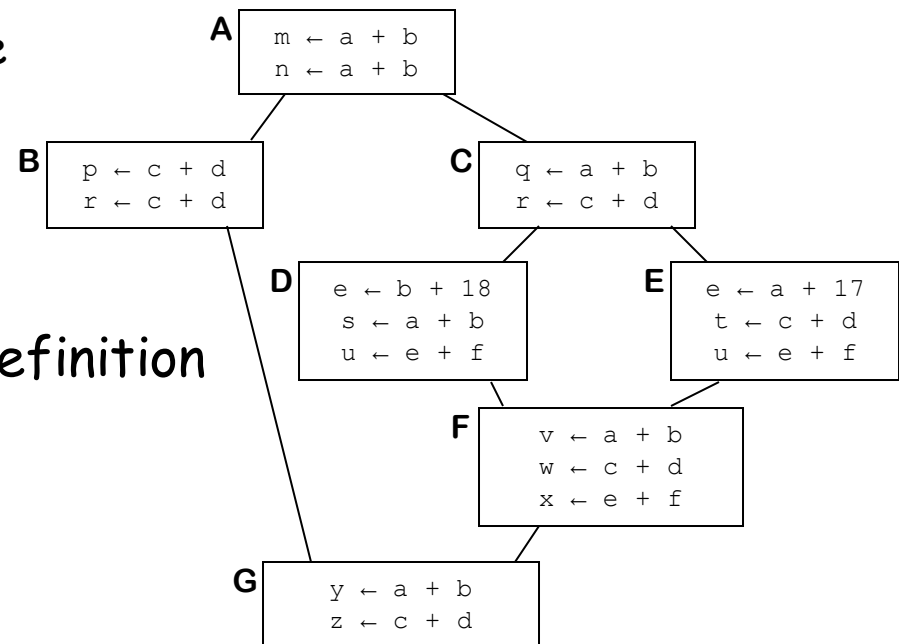
## Efficiency

- Use A's table to initialize tables for B & C
- To avoid duplication, use a scoped hash table
  - A, AB, A, AC, ACD, AC, ACE, F, G
- Need a VN  $\rightarrow$  name mapping to handle kills
  - Must restore map with scope
  - Adds complication, not cost

"kill" is a re-definition of some name

## To simplify matters

- Need *unique name* for each definition
- Makes name  $\rightarrow$  VN
- Use the SSA name space



The subscripted names from the earlier example are an instance of the SSA name space.

# SSA Name Space

(locally)



Example (from earlier):

## Original Code

$a_0 \leftarrow x_0 + y_0$   
★  $b_0 \leftarrow x_0 + y_0$   
 $a_1 \leftarrow 17$   
★  $c_0 \leftarrow x_0 + y_0$

## With VNs

$a_0^3 \leftarrow x_0^1 + y_0^2$   
★  $b_0^3 \leftarrow x_0^1 + y_0^2$   
 $a_1^4 \leftarrow 17$   
★  $c_0^3 \leftarrow x_0^1 + y_0^2$

## Rewritten

$a_0^3 \leftarrow x_0^1 + y_0^2$   
★  $b_0^3 \leftarrow a_0^3$   
 $a_1^4 \leftarrow 17$   
★  $c_0^3 \leftarrow a_0^3$

### Renaming:

- Give each value a unique name
- Makes it clear

### Notation:

- While complex, the meaning is clear

### Result:

- $a_0^3$  is available
- Rewriting just works



# SSA Name Space

(in general)

Two principles

- Each name is defined by exactly one operation
- Each operand refers to exactly one definition

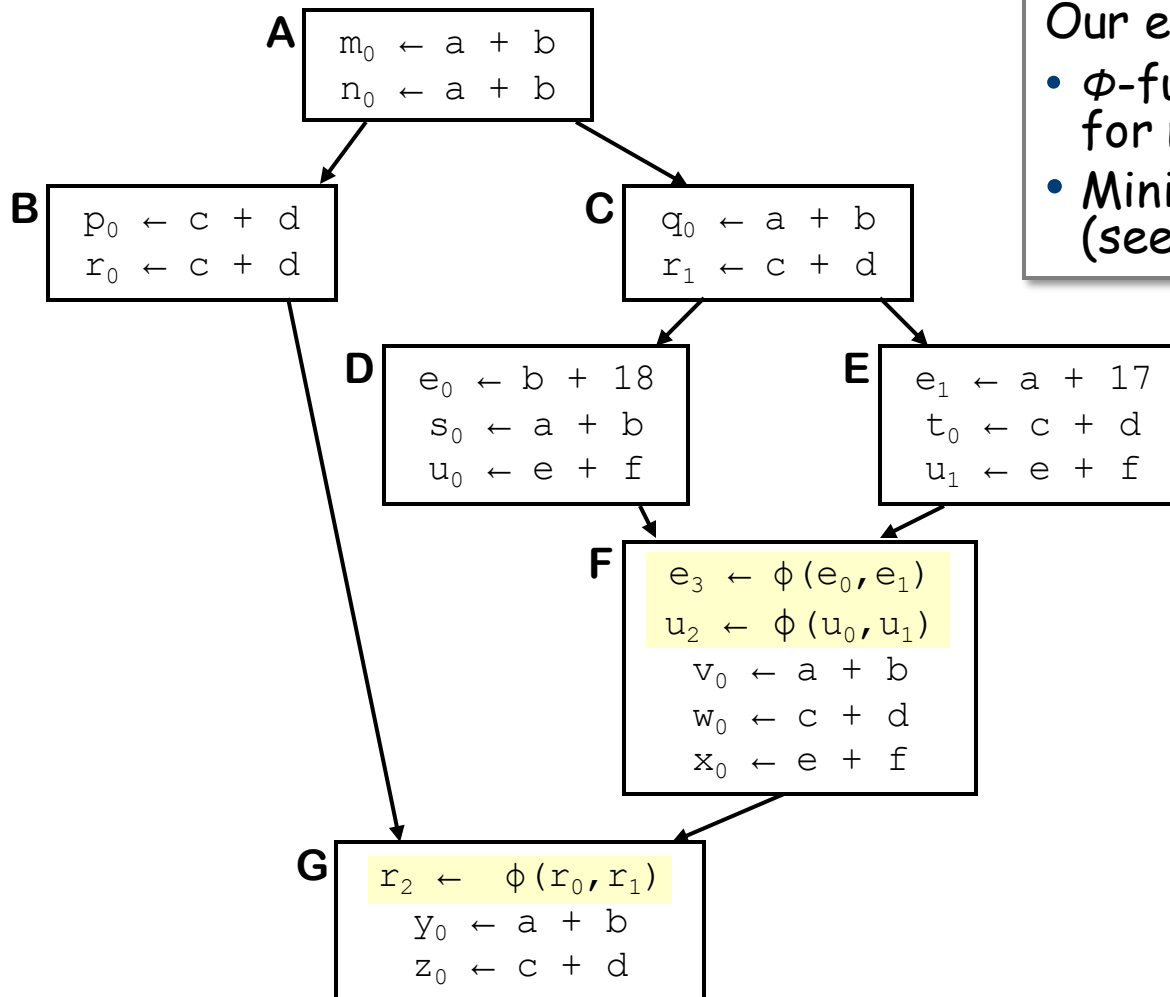
To reconcile these principles with real code

- Insert  $\phi$ -functions at merge points to reconcile name space
- Add subscripts to variable names for uniqueness





# Superlocal Value Numbering



Our example in SSA form

- $\phi$ -functions at join points for names that need them
- Minimal set of  $\phi$ -functions (see Chapter 9 in EaC)



# Superlocal Value Numbering

## The SVN Algorithm

*WorkList*  $\leftarrow$  { entry block }

Blocks to process

*Empty*  $\leftarrow$  new table

Table for base case

*while* (*WorkList* is not empty)

    remove a block *b* from *WorkList*

    SVN(*b*, *Empty*)

SVN( *Block*, *Table* )

*t*  $\leftarrow$  new table for *Block*, with *Table* linked as surrounding scope

    LVN( *Block*, *t* )

Use LVN for the work

    for each successor *s* of *Block*

        if *s* has just 1 predecessor

In the same EBB

            then SVN( *s*, *t* )

        else if *s* has not been processed

Starts a new EBB

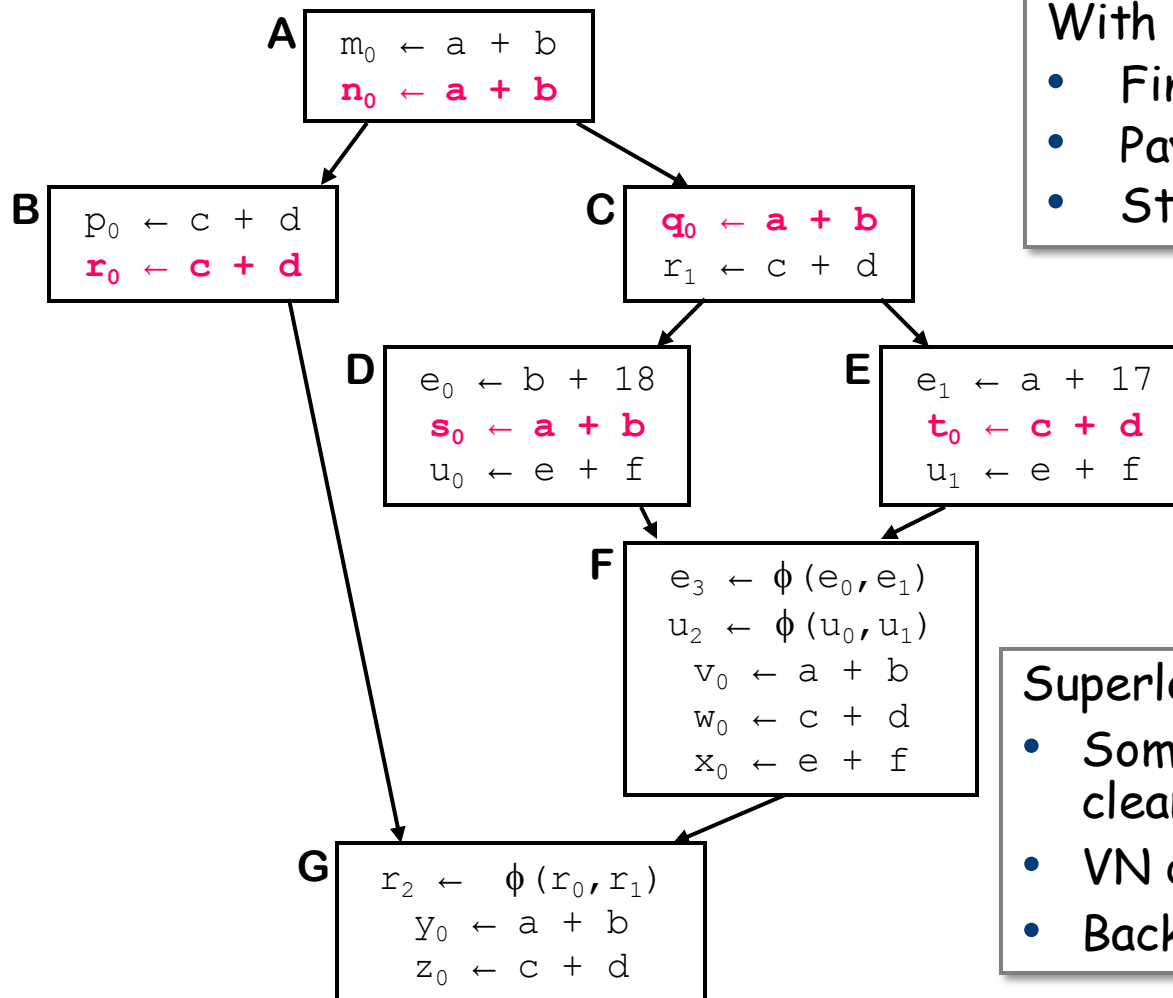
            then add *s* to *WorkList*

    deallocate *t*





# Superlocal Value Numbering



With all the bells & whistles

- Find more **redundancy**
- Pay minimal extra cost
- Still does nothing for F & G

Superlocal techniques


- Some local methods extend cleanly to superlocal scopes
- VN does not back up
- Backward motion causes probs



# Loop Unrolling

Applications spend a lot of time in loops

- We can reduce loop overhead by unrolling the loop

<pre>do i = 1 to 100 by 1   a(i) ← b(i) * c(i) end</pre>		<pre>a(1)  ← b(1) * c(1) a(2)  ← b(2) * c(2) a(2)  ← b(3) * c(3) ... a(100) ← b(100) * c(100)</pre>
	Complete unrolling	

- Eliminated additions, tests, and branches
  - *Can subject resulting code to strong local optimization!*
- Only works with fixed loop bounds & few iterations
- The principle, however, is sound
- Unrolling is always safe, as long as we get the bounds right



# Loop Unrolling

Unrolling by smaller factors can achieve much of the benefit

Example: unroll by 4

```
do i = 1 to 100 by 1  
  a(i) ← b(i) * c(i)  
end
```



```
do i = 1 to 100 by 4  
  a(i)    ← b(i) * c(i)  
  a(i+1) ← b(i+1) * c(i+1)  
  a(i+2) ← b(i+2) * c(i+2)  
  a(i+3) ← b(i+3) * c(i+3)  
end
```

Achieves much of the savings with lower code growth

- Reduces tests & branches by 25%
- LVN will eliminate duplicate adds and redundant expressions
- Less overhead per useful operation

But, it relied on knowledge of the loop bounds...



# Loop Unrolling

## Unrolling with unknown bounds

Need to generate guard loops

```
do i = 1 to n by 1
  a(i) ← b(i) * c(i)
end
```



Unroll by 4

```
i ← 1
do while (i+3 < n)
  a(i)    ← b(i) * c(i)
  a(i+1) ← b(i+1) * c(i+1)
  a(i+2) ← b(i+2) * c(i+2)
  a(i+3) ← b(i+3) * c(i+3)
  i ← i + 4
end
```

```
do while (i < n)
  a(i)    ← b(i) * c(i)
  i ← i + 1
end
```

Achieves most of the savings

- Reduces tests & branches by 25%
- LVN still works on loop body
- Guard loop takes some space

Can generalize to arbitrary upper & lower bounds, unroll factors



# Loop Unrolling

## One other unrolling trick

Eliminate copies at the end of a loop

```
t1 ← b(0)
do i = 1 to 100 by 1
  t2 ← b(i)
  a(i) ← a(i) + t1 + t2
  t1 ← t2
end
```



```
t1 ← b(0)
do i = 1 to 100 by 2
  t2 ← b(i)
  a(i) ← a(i) + t1 + t2
  t1 ← b(i+1)
  a(i+1) ← a(i+1) + t2 + t1
end
```

Unroll by LCM of copy-cycle lengths

- Eliminates the copies, which were a naming artifact
- Achieves some of the benefits of unrolling
  - Lower overhead, longer blocks for local optimization
- Situation occurs in more cases than you might suspect